

Estimating Statistical Power for Structural Equation Models in Developmental Cognitive
Science: A Tutorial in R

Elisa S. Buchberger¹, Chi T. Ngo¹, Aaron Peikert^{1, 2}, Andreas M. Brandmaier^{1,2,3}, & Markus
Werkle-Bergner¹

¹ Max Planck Institute for Human Development

² Max Planck UCL Centre for Computational Psychiatry and Ageing Research, Berlin,
Germany, and London, United Kingdom

³ Department of Psychology, MSB Medical School Berlin, Berlin, Germany

This work was conducted within the project “Lifespan Rhythms of Memory and Cognition (RHYME)” at the Max Planck Institute for Human Development. ESB is a fellow of the International Max Planck Research School on the Life Course (LIFE; <http://www.imprs-life.mpg.de/en>). AP is a fellow of the International Max Planck Research School on Computational Methods in Psychiatry and Ageing Research (COMP2PSYCH, <https://www.mps-ucl-centre.mpg.de/comp2psych>). MW-B received support from the Jacobs Foundation (Early Career Research Fellowship 2017–2019). CTN’s work is supported by a grant from the German Research Foundation (DFG; NG 191/2-1) and an Early Career Research Fellowship by the Jacobs Foundation (2021-1417-99). We are grateful to all members of the RHYME and LIME projects for valuable feedback on the project. Finally, we acknowledge support by the Max Planck Dahlem Campus of Cognition (MPDCC).

Correspondence concerning this article should be addressed to Elisa S. Buchberger, Lentzeallee 94, 14195 Berlin. E-mail: buchberger@mpib-berlin.mpg.de

Abstract

Determining the compositional structure and dimensionality of psychological constructs lies at the heart of many research questions in developmental science. Structural equation modelling (SEM) provides a versatile framework for formalizing and estimating the relationships among multiple latent constructs. While the flexibility of SEM can accommodate many complex assumptions on the underlying structure of psychological constructs, it makes a-priori estimation of statistical power and required sample size challenging. This difficulty is magnified when comparing non-nested SEMs, which prevents the use of the likelihood-ratio test. Sample size estimates for SEM model fit comparisons typically rely on generic rules of thumb. Such heuristics can be misleading because statistical power in SEM depends on a variety of model properties. Here, we demonstrate a Monte Carlo simulation approach for estimating a-priori statistical power for non-nested model comparisons in an SEM framework. We provide a step-by-step guide to this approach based on an example from our memory development research in children.

Keywords: structural equation modelling, statistical power, Monte Carlo simulation, sample size planning

Word count: NA

Estimating Statistical Power for Structural Equation Models in Developmental Cognitive Science: A Tutorial in R

1. Introduction

Over the past decades, many psychological constructs that had originally been conceptualized as unitary entities have been shown to be multifactorial phenomena. A prominent example are the early debates about component factors of intelligence (Cattell, 1971; Horn, 1970, 1978) or personality (McCrae & Costa, 1985). But also more specialized cognitive abilities such as executive functions (Hedden & Yoon, 2006; Miyake et al., 2000), attention (Mirsky, Anthony, Duncan, Ahearn, & Kellam, 1991) or memory (McClelland, McNaughton, & O'Reilly, 1995; Norman & O'Reilly, 2003; Schapiro, Turk-Browne, Botvinick, & Norman, 2017) are increasingly considered multi-process functions.

Understanding the underlying compositional structure of such multifaceted psychological constructs is challenging. The task becomes even more demanding when the development of component structures over time or the identification of (time-varying) differences between groups is the target of scientific investigations (Baltes, Reese, & Nesselroade, 1988). Structural equation modelling (SEM, Bollen, 1989; Jöreskog & Goldberger, 1975) provides a flexible framework that can accommodate a number of assumptions on the structure of key factors underlying a given construct and their inter-dependence. Indeed, SEM has become an increasingly popular tool in psychological science over the past decades (for a review, see MacCallum & Austin, 2000).

Since Cohen's classic study on statistical power (Cohen, 1962), numerous researchers from psychological and neurocognitive science have stressed the issue of insufficient statistical power and its consequences on the interpretation of scientific findings in the face of inadequate sample sizes Vankov, Bowers, & Munafò (2014). However, determining required sample size in an SEM approach is far from trivial. While the investigation of target effects in SEM can be achieved via analytical computations (Satorra & Saris, 1985),

tackling more complex research questions that rely on the comparison of non-nested models poses a challenge in determining sufficient sample sizes.

In this study, we implemented a Monte Carlo based simulation approach and provide a step-by-step guide for conducting randomization-based analyses for a-priori sample size estimations for non-nested model comparison. Specifically, we exemplify this approach with a specific research question, focusing on the componential structure of memory processes in early childhood (for theoretical background see Buchberger, Brandmaier, Lindenberger, Werkle-Bergner, & Ngo, 2022). Memory developmental research represents an excellent example for this methodological approach, as researchers from different fields have put forth competing ideas on the underlying structure of memory in childhood: While a rich body of empirical work from many decades has focused on the dichotomy between episodic and semantic memory (Tulving, 1972), more recent neurocomputational theories have introduced a process-focused approach to memory development. Tenets from contemporary memory models invite the hypothesis that three kinds of neurocomputations may complementarily support adaptive memory functioning, thereby assuming a tri-partite structure of memory (McClelland et al., 1995; Norman & O'Reilly, 2003; Schapiro et al., 2017). Such competing theories on the underlying structure of memory can be addressed via model comparisons in an SEM framework. In this tutorial, we address two overarching issues that affect estimates of statistical power for such comparisons: (i) the separability of the theoretical constructs and (ii) the reliability with which the constructs are being measured. We employ this example to demonstrate the procedure and the utility of such methods, with the aim of illustrating a transferable approach to other research questions in different domains and disciplines.

1.1 Structural Equation Modelling in Developmental Cognitive Neuroscience

Structural Equation Modeling (SEM) uses information about the relationship between multiple measured variables to uncover the structure of unobservable constructs, rendering this methodological framework a powerful tool in the field of developmental cognitive

neuroscience (e.g., Baltes et al., 1988; Kievit et al., 2018). Most notably, it enables the translation from theories to testable mathematical models that can simultaneously include numerous observable indicators and multiple latent constructs of interest (such as behavioral constructs and neural correlates). Thus, it allows simultaneous modelling of the relationships among theoretical constructs and their co-development across different developmental windows.

In an SEM framework, theoretical constructs of interest are operationalized as latent variables that capture the underlying commonalities across a set of (measured) manifest variables. Manifest variables can directly be observed (such as items in a questionnaire or indices derived from a behavioral task) and are commonly illustrated as squares in graphical depictions of SEM. Latent variables represent hypothetical constructs that are not directly observable (such as a latent factor of intelligence or memory), but are inferred from a selection of manifest variables. In graphical depictions of SEM, latent variables are commonly shown as circles. The strength of a loading from a manifest onto a latent variable represents the extent to which the observed variance is accounted for by the latent factor. In other words, the loading indicates how well a given manifest variable captures the latent construct.

The relationship among manifest and latent variables can be described in a mathematical model, which specifies all assumed parameters. Two models M1 and M2 are referred to as nested if the parameter space in a more restrictive model M2 represents a subspace of the parameter space of the more general model M1 (Bentler & Bonett, 1980). This means that the two models only differ with regards to the specification of one or multiple parameters. Nesting of models is usually achieved via constraining free parameters from M1 to equality or to known constants. The concept of nesting represents a convenient characteristic when comparing competing models, as it allows the use of a Likelihood Ratio test to evaluate their relative model fits (Satorra & Saris, 1985). When comparing non-nested models, researchers usually have to revert to model fit indices that indicate how

well a given model describes the data, such as the Bayesian Information Criterion (BIC, Raftery, 1986; Schwarz, 1978), the Akaike Information Criterion (AIC, Akaike, 1974), the Root Mean Square Error of Approximation (RMSEA, Steiger, 2016), and the Comparative Fit Index (CFI, Bentler, 1990). Such indices can guide heuristic decisions that do not hold any statistical guarantees (e.g., a pre-specified Type-I error level).

1.2 Statistical Power in SEM

Over the past decades, the issue of low statistical power and its consequences for the interpretation of scientific findings has gained awareness in neurocognitive and psychological research (Cohen, 1988; MacCallum & Austin, 2000; Maxwell, 2004; Rossi, 1990). Statistical power refers to the probability of rejecting a null hypothesis, when it is indeed false (i.e. the probability of not committing a Type-II error) and is directly linked to the sample size of a study, the magnitude of the targeted effect, and the reliability of measurement (Brandmaier, Oertzen, Ghisletta, Lindenberger, & Hertzog, 2018; Cohen, 1988). While low statistical power decreases the scientific utility of any given study, it also decreases the likelihood that a given significant result actually reflects a true effect (Button et al., 2013). To allow well designed empirical research to detect effects in the sample under investigation, a-priori estimates of statistical power (that is computation of power estimates before conducting the study) and thus informed decisions on required sample size are crucial to prevent underpowered research (Button et al., 2013). While the desired level of statistical power can depend on specific aspects of the research question at hand, a typical convention for adequate power in the behavioral sciences is 0.8 (Cohen, 1988). However, others have argued that there is no reason to prefer Type-I errors over Type-II errors and thus one should better aim for a statistical power of 0.95, if the level of significance is kept at the conventional 5%. Note that of course, that considerations on statistical power and thus sample size estimates need to be balanced with the probability of committing Type-I errors, that is the probability to reject a null hypothesis if it is indeed true.

Various rules of thumb on required sample size in SEM have been recommended, including setting an absolute minimum of observations across the board (Boomsma, 1985) and adjusting to the model complexity (i.e., setting a number of observations per estimated parameter, Bentler & Chou, 1987). Unfortunately, such heuristics can be misleading, as statistical power in SEM is heavily influenced by parameters beyond the number of indicators per latent construct, e.g. the magnitude of factor loadings (Wolf, Harrington, Clark, & Miller, 2013). Further, determining the required sample size for a given study depends on the research question. One common utility of SEM is to test whether a given effect in a model exceeds a specific threshold, such as testing whether the correlation of latent factors exceeds zero, or whether a specific parameter estimate differs between groups (i.e., target effect). Another common utility is to determine how well a given model describes the data, and/or whether one model describes the data better than a competing model (i.e., model comparison). These different research questions require different types of power: the power to detect a target effect vs. the power to detect model misspecification (Wang & Rhemtulla, 2021). Further, theoretical assumptions on competing models can lead to the necessity to compare non-nested models, which prevents the use of regular χ^2 -based power estimates and therefore requires alternative approaches to determining statistical power for SEM.

1.2.1 Statistical Power to Detect a Specific Effect of Interest. Central to many studies that employ SEM is the question of whether a specific parameter in a model (e.g., strength of a specific correlation or the slope parameter in a latent regression model) is different from a given value (e.g. Canada, Hancock, & Riggins, 2021). Imagine that a researcher aims to test whether the correlation between two latent constructs significantly differs from zero. To this end, they would compare the model fit between one model in which the correlation parameter is freely estimated and another model in which it is fixed to zero (Satorra & Saris, 1985). The corresponding null hypothesis in this case states that the parameter restrictions hold in the population. The difference in model fit will follow a χ^2 -distribution with degrees-of-freedom (df) equal to the difference of freely estimated

parameters between the two models if the null hypothesis is true (Neale, 2000, in this example $df = 1$, because only the correlation between the two latent factors is fixed in the restricted model). If restricting the parameter of interest results in a significantly poorer model fit, this would suggest that the respective parameter indeed significantly differs from zero. Following the logic outlined above, researchers can determine the statistical power for detecting a target effect in an SEM framework before conducting the study. A key aspect for this a-priori power calculation hinges on translating the differences in a specific parameter estimate into an effect size. Such translation can be achieved by investigating the discrepancy between the model-implied variance-covariance matrices associated with (1) the population model including the “true” parameter values and (2) the hypothesized model. The discrepancy between both matrices is quantified based on a specific fit function (for details on fit functions see Bollen, 1989). Recently, several user friendly tools have emerged that can guide modelers in quantifying effect sizes of a target effect to estimate statistical power in SEM (e.g. the R packages `semPower` Jobst, Bader, & Moshagen, 2021; or `power4SEM` Jak, Jorgensen, Verdam, Oort, & Elffers, 2020 for analytical approaches; and the Shiny app `pwrSEM` Wang & Rhemtulla, 2021; the interactive study planner tool `LIFESPAN` Brandmaier, Oertzen, Ghisletta, Hertzog, & Lindenberger, 2015; or the R package `simsem` Pornprasertmanit, Miller, Schoemann, & Jorgensen, 2021 for simulation-based approaches). As the investigation of a target effect practically translates into the comparison of nested models, researchers can in these cases analytically determine the statistical power and use this information for decisions on required sample size (Satorra & Saris, 1985).

1.2.2 Statistical Power for Model Comparison. A second question that is of interest for many researchers – especially in developmental cognitive neuroscience – pertains to identifying the one theoretical model (from a set of competing models) that best explains the given data (Henson et al., 2016; Miller, Giesbrecht, Müller, McInerney, & Kerns, 2012; Nyberg, 1994). In the case of competing models that are nested, the analytical approach to this question (and therefore also considerations on statistical power) can be addressed

analogously to the procedure outlined above (see Jobst et al., 2021 for a step-by step tutorial). However, comparing non-nested models prohibits the use of regular χ^2 -based statistics and therefore poses additional methodological challenges in determining the ‘best’ model. Nevertheless, researchers should strive for a methodological approach that matches the theoretical assumptions, rather than vice versa, that is, moving away from well-grounded theoretical considerations to meet methodological constraints. In cases where analytical solutions to questions of statistical properties like statistical power are not available, randomization-based techniques offer an excellent alternative (Efron & Tibshirani, 1994). In particular, Monte Carlo simulations have proven useful to bridging this gap (Muthén & Muthén, 2002, see section 2.1 for more details on Monte Carlo simulations).

1.3 Example: Competing Models on the Compositional Structure of Memory

For many decades, the compositional nature of memory in adults as well as its maturation across development has been of great interest in psychology, cognitive science, neuroscience and artificial intelligence research. Different models on the compositional structure of memory have been heavily debated (McClelland et al., 1995; Norman & O’Reilly, 2003; Schapiro et al., 2017; Tulving, 1972).

The most simplistic characterization of memory structure is to assume a unitary ability underlying different memory demands, akin to a g-factor of memory (Spearman, 1904). Such a unitary model of memory suggests no differentiation of component processes within declarative memory, but rather claims that the performance on various types of memory demands is grounded in a single ability factor.

An alternative prominent view is the classic dichotomy between episodic and semantic memory systems (Squire, 1987; Tulving, 1972). Such a bi-partite architecture of memory postulates a division between one component responsible for learning specific events embedded in their temporal and spatial context (episodic memory) and a second component

responsible for learning regularities and extracting generalized knowledge (semantic memory). Previous research comparing a unitary vs. a bi-partite account of memory in an SEM framework has supported the idea of separable memory factors underlying declarative memory, at least in adults (Nyberg, 1994). Adopting a process-oriented view, recent computational models of memory have argued for a functionalist distinction between memory specificity and generalization through a labor division between the hippocampus and the cortex (McClelland et al., 1995; Norman & O'Reilly, 2003). Here, a set of neurocomputations support different memory demands. Important to learning specific episodes are pattern separation that discriminates between similar memories through the reduction of representational similarity, and pattern completion that retrieves linked associations among co-occurring elements (Marr, 1971; Norman & O'Reilly, 2003; Rolls, 2016). Both of these computations are thought to be specialties of the hippocampus. In contrast, the cortex is well suited to slowly learn statistical regularities that enables generalization (McClelland et al., 1995; Norman & O'Reilly, 2003). Interestingly, rapid generalization also relies on the hippocampus, either via retrieval mechanisms of related episodes (Kumaran & McClelland, 2012) or via a distributed coding scheme carried by a specific subset of the hippocampal circuitry (Schapiro et al., 2017). From this vantage point, memory abilities take shape of a tri-partite structure, encompassing pattern separation, pattern completion and generalization as three separable mnemonic processes.

Which of these views best explain memory abilities across early development? This question requires adequate statistical methods that can adjudicate among multiple competing ideas on the underlying structure of memory. The simulation in the following section of this paper will address methodological challenges in study planning when aiming to compare three hypothetical models: a unitary model (Model 1), a bi-partite model (Model 2) and a tri-partite model (Model 3) on the compositional structure of memory (Figure 1). Specifically, we will focus on the issue of estimating statistical power for this model comparison. Here, we address the impact of two main aspects in particular: (i) the

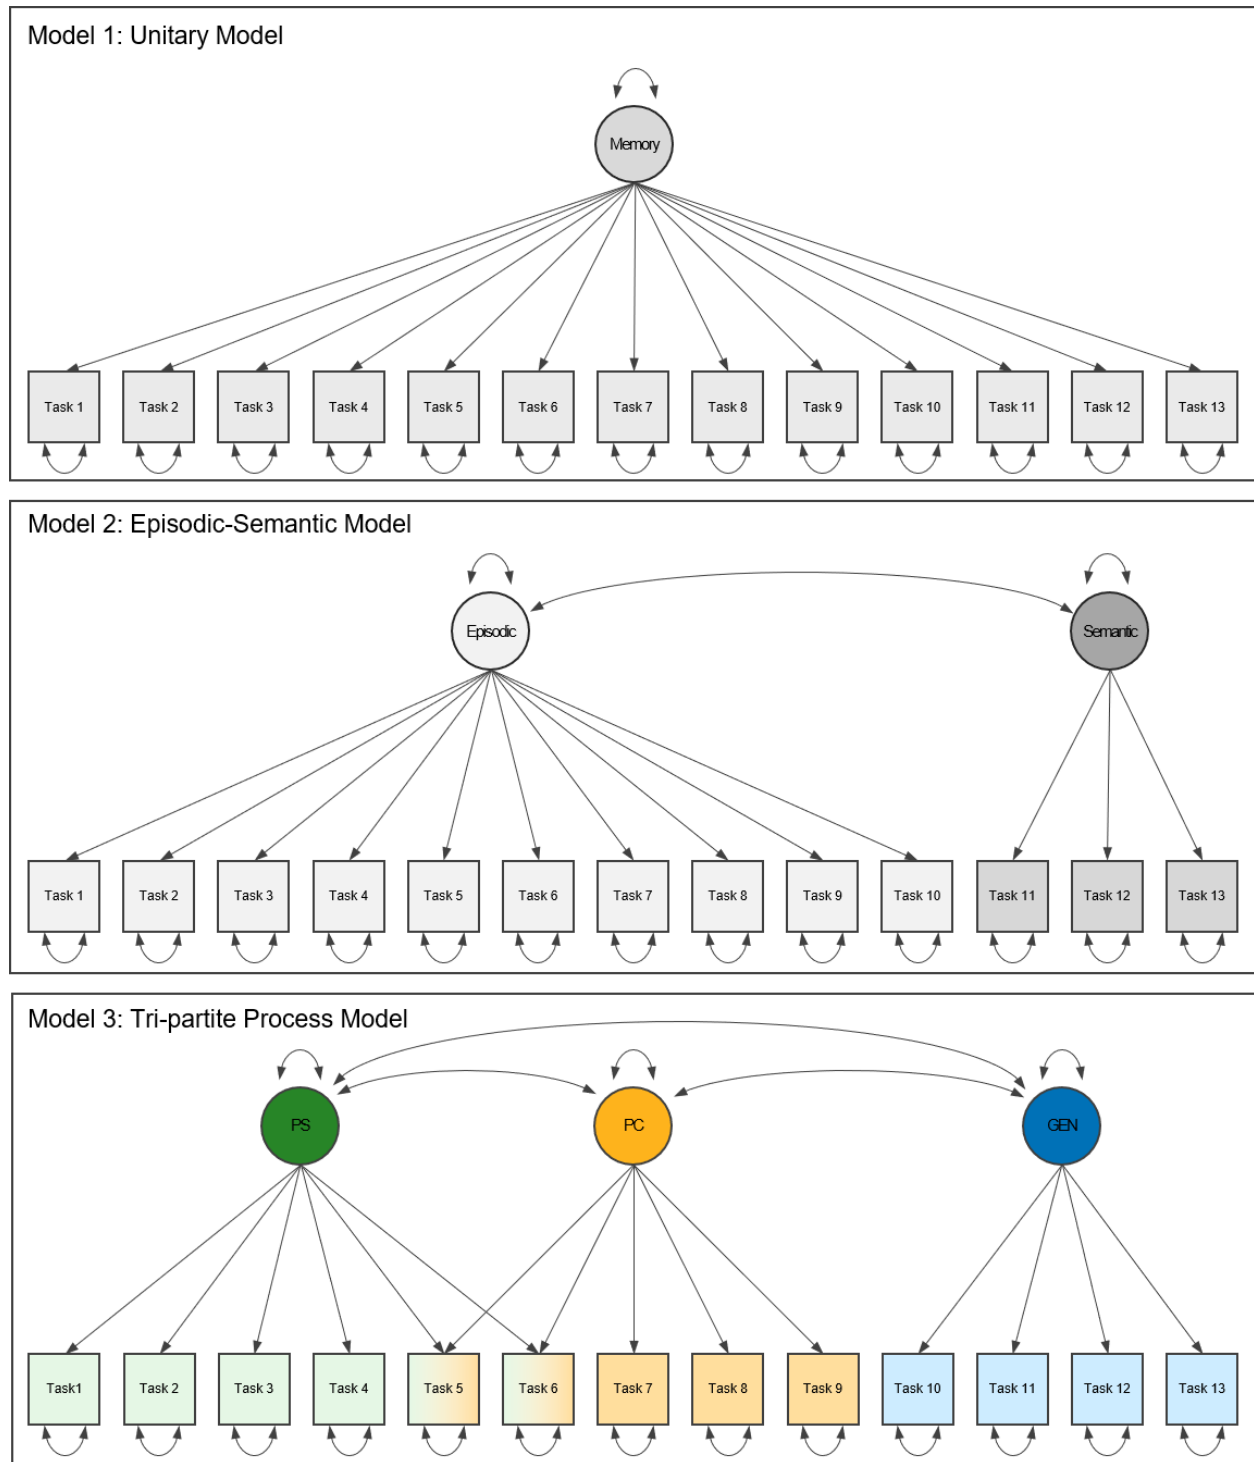


Figure 1. Path diagrams of three competing SEMs on the structure of memory. Squares represent observed performance measures on the behavioral tasks. Circles represent latent constructs (PS = Pattern Separation, PC = Pattern Completion, GEN = Generalization). Single-headed arrows indicate regressions and double-headed arrows indicate covariances. See Buchberger et al., (2022) for in depth description of the models.

separability of the theoretical constructs, i.e. inter-relatedness of the latent factors in the competing models and (ii) the reliability of the measures, i.e. the loadings of the manifest variables. Details on the theoretical bases and the selection of methodological indicators for each model can be found in Buchberger et al. (2022). Important for this context is the fact that our specification of Models 2 and 3 results in a non-nested model comparison (due to the cross-loadings for the latent factors pattern separation and pattern completion and different allocation of the indicator Task 10). Therefore a-priori estimations of statistical power call for a simulation-based approach.

2. Methods

In the following, we demonstrate how to conduct a simulation-based power analysis for non-nested model comparisons. For the sake of illustration, we will refer to the example on the compositional structure of memory outlined in section 1.3 throughout. For this example, we aim to derive estimates of required sample size to identify the tri-partite model as the best fitting model, if it indeed generated the data. In a step-by-step guide, we therefore show how to specify the competing models in R, how to set up different design conditions for the simulation, how to define the functions to generate, analyze and summarize synthetic data and how to execute the simulation. Based on two separate simulations, we show how the separability of theoretical constructs (section 3.1) and the reliability of the manifest variables (section 3.2) impacts estimates of statistical power for model comparisons by varying inter-factor correlations and factor loadings across different iterations of the simulation. Finally, we employ the same simulation-based approach to investigate Type-I error rates, that is erroneously identifying the tri-partite model as the best fitting model if the bi-partite or unitary model generated the data (section 3.3).

2.1 How does a simulation work: Power-estimations for the comparison of non-nested SEMs

To investigate statistical power for deciding between competing theoretical non-nested models, we suggest a Monte Carlo simulation-based approach (Metropolis, Rosenbluth, Rosenbluth, Teller, & Teller, 1953). Monte Carlo studies represent a computer-intensive simulation, that allows to approximate statistical power for a given study design (Muthén & Muthén, 2002). While Monte Carlo simulations can be implemented for a variety of research questions, the flexibility of the approach makes them especially well-suited for obtaining a-priori power estimates for the comparison of non-nested SEMs. Furthermore, analytical approaches are tied to assumptions (e.g., no missing data, multivariate normality), whereas simulation-based approaches allow for arbitrary data generating processes. In Monte Carlo studies, synthetic data is repeatedly generated with a set of different hypothesized parameter values and analyzed across all samples. Summary statistics from the entirety of the simulated datasets are then used to draw conclusions (Muthén & Muthén, 2002; Paxton, Curran, Bollen, Kirby, & Chen, 2001). The general procedure of Monte Carlo simulations follows a common core structure: (i) generate – in which multiple synthetic datasets are generated based on the hypothesized models, (ii) analyze – in which the synthetic data are analyzed for every iteration of the simulation, (iii) summarize – in which the results are pooled over all simulation iterations. This general procedure is applicable to a variety of scenarios. When comparing competing SEMs, it is necessary to specify the competing theoretical models, and decide on the simulation design, that is which model parameters of interest should be modified across the simulations. Finally, researchers need to evaluate the results from the simulation in order to derive practical implications from the simulation (see Figure 2). While these steps can be manually coded in a statistical research software, such as R, the implementation of a simulation can become increasingly complex and error-prone with increasing number of conditions that are being simulated. A useful tool which can guide novice simulators in setting up a Monte Carlo simulation is the R package `simDesign`

(Chalmers & Adkins, 2020), which facilitates the implementation of the internal logic of generate – analyze – simulate and can accommodate various research designs. All simulation steps in this study were therefore implemented within the `simDesign` package.



Figure 2. Schematic overview of the simulation process

2.1.1 Step 1: Specifying Competing SEMs. The first step of the simulation is translating the theoretical assumptions about latent constructs of interest and the manifest variables into an SEM. This step entails a) specifying the number of latent factors (how many theoretical constructs are thought to influence the observed behavior in the manifest variables?), and b) the allocation of each manifest variable to (at least one) latent variable (which of the manifest variables captures the respective latent construct?). In R, the structure of a SEM can be easily implemented using `lavaan` syntax (Rosseel, 2012). Typically, useful operators for specifying a model entail factor loadings (`=~`), (co-)variances (`~~`), and means or intercepts (`~1`). The tri-partite model (Figure 1, Model 3) could be specified in `lavaan` syntax as the following:

```
model3 <-
  'ps  =~ task1 + task2 + task3 + task4 + task5 + task6
    pc  =~ task5 + task6 + task7 + task8 + task9
    gen =~ task10+ task11 + task12 + task13'
```

The model is defined by a model string, which specifies three latent factors of memory (Pattern Separation [PS]; Pattern Completion, [PC]; Generalization, [GEN]) which are

measured through the manifest variables (Tasks 1 to 13). Note that lavaan entails several user-friendly functions for fitting models that automatically include default settings for model components that are not explicitly specified (e.g. adding residual variances, covariances of exogenous latent variables). Thus, the model syntax can be kept very concise (check documentation of the respective function for full information on default settings).

Estimating statistical power for model comparisons requires that researchers specify each of the competing models and all aspects in which they differ from one another (e.g., in number of latent factors, allocations of manifest variables; see supplementary R code on GitHub for specifications of Models 1 and 2).

2.1.2 Step 2: Simulation Design. When defining the structure of the theoretical model for data generation, it is crucial to specify all parameters in the model, particularly the loading strength of the manifest variables, along with the means and (co-)variances of the latent factors. This step also requires the specification of which parameter values should be kept constant or modified across different simulation iterations. Parameters for which there are strong assumptions (e.g. derived from previous work) can be entered in the models and kept constant. Parameters for which the impact on statistical power is of interest should be included as a condition in the simulation design. For example, if the question pertains to how the reliability of the measures may impact the estimates of required sample size, different levels of loading strength can be included in the conditions for which data is being simulated. The `createDesign()` function of the `simDesign` package can be used to create a full list of conditions for the simulation. All parameters that should be varied in the simulation are included as arguments in the `createDesign()` function. Full crossings of all possible values for the different parameters are then generated as conditions for the simulation. Note, however, that the number of simulation conditions and therefore the required computation time increases exponentially for full combinations of all simulation conditions. Researchers should thus be mindful about the number of conditions in a given simulation to avoid overly computationally intensive simulations, which further can be increasingly difficult to interpret.

If only specific combinations of the parameters in the simulation are of interest for a given research question, the resulting design data frame can be adjusted accordingly, e.g., by deleting specific rows from the data frame. Below is our design data frame that includes different simulation conditions for (i) the loading strength of the manifest variables and (ii) the sample sizes:

```
design_load_model3 <- createDesign(
  loading_strength = seq(0.5, 0.9, 0.1),
  sample_size = seq(50, 200, 5),
  # [...]
)
```

Here, the design for the Monte Carlo simulations consists of five loading strength conditions (i.e., 0.5 to 0.9 in increments of 0.1). The sample size is varied from 50 to 200 in increments of 5. In total, data for 155 conditions (5×31) are generated.

2.1.3 Step 3: Generate. Next, we define the mechanism to generate the synthetic data. An easy way to generating synthetic data within an SEM framework in R is the `simulateData()` function available in the `lavaan` package. Within the `generate` function, all parameters that are predetermined across iterations can be directly entered into the generating model syntax. However, all parameters that are supposed to be varied between different iterations of the simulation need to be accessed through the `condition` argument that is returned through the `createDesign()` function in step 2. Given our interests in the effect of loading strength and strength of covariance between the latent factors, we define these parameters in step 2 (Design) such that they can now be called via the `condition` argument in the `generate` function:

```
generate_data <- function(condition, fixed_objects = NULL) {
  a <- condition$loading_strength
  # [...]
}
```

```

b <- a / sqrt(2 + 2 * condition$cov1)

pop_model <-
  paste0(
    "ps =~" ,
    paste0(a, "*task", 1:4 , collapse = "+"),
    "+",
    b,
    "*task5 +",
    b,
    "*task6;",
    "pc =~",
    b,
    "*task5 +",
    b,
    "*task6 +",
    paste0(a, "*task", 7:9 , collapse = "+"),
    ";",
    "gen=~",
    paste0(a, "*task", 10:13 , collapse = "+"),
    ";",
    paste0("task", 1:13, " ~~ ", 1 - a ^ 2 , " *task", 1:13, ";" ,
           collapse = ""),
    "ps ~~",
    condition$cov1,
    "*pc;",
    "ps ~~",
    condition$cov2,

```

```

    "*gen;",
    "pc ~~",
    condition$cov3,
    "*gen"
  )
dat <-
  data.frame(simulateData(pop_model, sample.nobs = condition$sample_size))
}

```

Here, we define the generate function for the tripartite model. The factor loadings specified in step 2 (Design) are called from the condition argument such that they are varied across different iterations of the simulation, and are defined as a for syntax readability. For standardization purposes (that is, the total variance of each manifest variable summing up to 1), the loadings of indicators Task5 and Task6 are set to $\frac{a}{\sqrt{(2+2 \cdot cov1)}}$ to account for cross-loadings on multiple factors and are defined as b to ease syntax readability ($cov1$ = covariance between the latent factors pattern separation and pattern completion). The population model is defined by a model string, which specifies three latent factors (ps, pc, gen), measured by the 13 manifest indicators (task1 – task13). The residual variances of the manifest variables are set to $1 - a^2$ for standardization purposes. The covariance between the latent factors is set to the respective value from the condition argument. The data are simulated from the model syntax via the `simulateData()` function from the `lavaan` package and stored as `dat`.

Note that the Monte Carlo simulation method does not limit researchers to the simulation of normally distributed data, but can be used with any form of data generating mechanism. Further, the simulation-based approach allows researchers to anticipate and investigate the effects of (planned) missingness in the data on estimates for statistical power (Brandmaier, 2020; Schoemann, Miller, Pornprasertmanit, & Wu, 2014). Here, different

amounts and patterns of missingness can be entered into the simulation by removing data from the generated datasets. The resulting data can then be analyzed and summarized to investigate the effects of attrition, planned or random missingness on the estimates for statistical power. The implementation of these characteristics of the data is beyond the scope of this tutorial. However, in research scenarios in which violations of the normality assumption or missing data can be expected, researchers should consider a simulation-based approach to estimate statistical power, as the approach presented here is free from any assumptions on the data distribution.

2.1.4 Step 4: Analyze. In this step, we define the analyze function, which specifies how the data generated in step 3 should be analyzed. For the comparison of non-nested models, we rely on a combination of relative and absolute model fit parameters. We fit each of the competing models to each of the simulated data sets using the `sem()` function from the `lavaan` package and extract different measures of model fit using the helper function `get_fitmeasures()` (see `funs.R` in the supplementary material on GitHub), namely the Comparative Fit Index (CFI, Bentler, 1990), root mean square error of approximation (RMSEA, Steiger, 2016) and the Bayesian Information Criterion (BIC, Raftery, 1986). The function returns the fit measures, the simulation model that was used in the respective iteration and information on model convergence. The analyze function for the example on memory composition could thus look like this:

```
analyze_results <-  
  function(condition, dat, fixed_objects = fixed_objects) {  
    fits <- map(fixed_objects, ~ sem(., dat, std.lv = TRUE))  
    ms <- map(fits, get_fitmeasures)  
    ret <- list(  
      fitmeasures = ms,  
      sim_model = condition$sim_model,  
      sample_size = condition$sample_size,
```

```

    loadingstrength = condition$loading_strength,
    covariance = condition$cov1,
    converged = map_lgl(fits, is_converged)
  )
}

```

2.1.5 Step 5: Summarize. In step 5, we define how the results from the analyze step should be summarized. To this end, we extract how often the model used to generate the data (step 2) was actually recovered as the best fitting model. We implemented the following decision rules for identifying the best fitting model: (i) We define cut-off criteria for an adequate model fit according to established rules as $CFI > .95$ and $RMSEA < 0.06$ (Hu & Bentler, 1999). (ii) Among the models that meet the cut-off criteria, we select the best fitting model for each data set based on the relative model fit (lowest BIC value). The goal of (i) and (ii) is to check whether the ground truth model that generated the data can be identified as the best fitting model in a given simulation run. For cases in which a given dataset does not converge for one of the models, we consider the outcome of these iterations as failing to recover the correct model. We summarize the proportion of successful model recovery for each combination of model parameters in the simulation separately. The summarize function that concatenates the results across simulation iterations for the example on the structure of memory in childhood could thus look like this:

```

summarize_results <-
  function(condition, results) {
    rmsea_cut = 0.06
    cfi_cut = 0.95
    tidied <- results %>%
      bind_rows(.id = "rep") %>%
      mutate(model = names(fitmeasures)) %>%

```

```

    unnest_wider(fitmeasures) %>%

    mutate(admissable = (rmsea < rmsea_cut) & (cfi > cfi_cut))

ret <-

tidied %>%

  group_by(rep) %>%

  summarise(

    sim_model = unique(sim_model),

    best = if (any(converged == FALSE)) {

      0

    } else if (is_empty(bic[admissable])) {

      0

    } else {

      parse_number(model[bic == min(bic[admissable])])

    }

  ) %>%

  mutate(

    model3_chosen = best == 3,

    sim_model_chosen = best == sim_model

  ) %>%

  with(c(

    best_3 = mean(model3_chosen),

    best_sim = mean(sim_model_chosen)

  ))

  return(ret)
}

```

First, the function restructures the results from the analyze function to obtain a tidy

data frame with the fit measures for each model in a separate row. This data is then summarized for each iteration. The function identifies a) the ground truth model that generated the data in this simulation iteration (`sim_model`) and b) the model which produced the best model fit (`best`). Here, the best model is defined as “0” if one of the three theoretical models did not converge, or if none of the models met the cut-off criteria as defined above. Otherwise, the best model is selected among the ones that met the cut-off criteria based on the relative model-fit. The function returns the argument `best_3`, indicating how often across all iterations Model 3 was identified as the best fitting model, and `best_sim`, indicating how often the model that simulated the data was identified as the best fitting model (note that in the current analysis, both values are identical. `best_sim` can however be useful to investigate Type-I error, see section 3.3).

2.1.6. Step 6: Running and Evaluating the Simulation. In a final step, we call the `runSimulation()` function from the `simDesign` package to execute the consecutive steps outlined above. Running the simulation for our example could look like this:

```
results_cov_model3 <-  
  runSimulation(  
    design = design_cov_model3,  
    replications = 10000,  
    generate = generate_data,  
    analyse = analyze_results,  
    summarise = summarize_results,  
    store_results = TRUE,  
    parallel = TRUE,  
    packages = c('SimDesign', 'lavaan', 'tidyverse'),  
    fixed_objects = fixed_objects  
  )
```

In this function, the pre-defined functions `generate_data()`, `analyze_results()` and `summarise_results()` are called. The number of replications is specified in the `replications` argument. To ensure a reasonable stability of the results, a sufficient number of replications is recommended. In our example, we simulated 10,000 datasets per condition. This leads to a standard error of the simulated power of less than 0.5%. The `store_results` argument can be used to store the individual results generated by the `summarize` function, allowing users to investigate the outcome on every iteration (in this example, the data frame containing all fit measures and convergence information for each iteration). Note however, that storing the individual results increases computation time and size of the resulting outcome arguments. The argument `parallel` can be used for parallel processing across multiple cores (e.g. when running the analysis on a computer cluster) to reduce computation time. If not specified otherwise, all available cores will be used per default when `parallel = TRUE`.

The output can then be used to examine power curves (see Figure 3) for each of the parameters of interest in the simulation. Power curves are line plots that depict how changing one variable (e.g. sample size) would affect the power of the test. In the example, we can plot power curves demonstrating the impact of the separability of the theoretical constructs, implemented as the strength of the covariance between the latent factors, or the impact of the reliability of the measures, implemented as the strength of the loadings, on power estimation. Importantly, these data reveal the relative importance of different model parameters on sample size requirements. To obtain easily readable outputs, we plot the sample size from the simulation (x-axis) against the frequency with which the data-generating model was correctly recovered as the best fitting model (y-axis). To derive estimates on the required sample size in the example presented here, we aimed for a power of 95% in recovering the ground truth model as the best fitting model.

3. Results: Power Estimates for Investigating the Compositional Structure of Memory

3.1 What is the Impact of the Separability of Theoretical Models on Estimates of Statistical Power?

In the example of comparing competing models of memory structure, we first examined the separability of the theoretical models by investigating the influence of the latent factors' correlations on statistical power estimates. In our example simulation, we investigated the statistical power to retrieve Model 3 as the best fitting model, when it indeed generated the data, while varying sample sizes from 50 to 200 and factor inter-correlations from 0.1 to 0.9. In this case, we restricted the simulated correlations between the latent factors to positive values as there is no indication in the literature to suspect a negative relation among these memory capacities (e.g., Tucker-Drob, Brandmaier, & Lindenberger, 2019). Figure 3A displays the obtained power curves illustrating the estimated statistical power dependent on the total sample size and the factor inter-correlation. For this estimate, the loading of the manifest variables was fixed at 0.7 and the inter-correlation between all latent factors were varied simultaneously.

As expected, statistical power increased with increasing sample size. However, the separability of the latent factors mattered: The higher the correlations between factors, the less able we were to distinguish multi-factor models from a one-factor model. That is, the smaller the difference between competing models, the greater the sample size that would be required to dissociate between them with adequate power. For very high levels of inter-correlation (0.8 and higher) the increase of statistical power with increasing sample size was remarkably lower than for lower levels of inter-correlation (0.7 and below). When the inter-factor correlation was 0.9, not even a larger sample size ($N=200$) achieved an adequate level of power of 80%. These results are not surprising given that increasing the inter-correlations between the latent factors renders the models more similar to each other, and thus makes their distinction more challenging. For moderate levels of inter-correlation

(≤ 0.7) a statistical power of 95% was reached for $N = 120$.

3.2 What is the Impact of the Reliability of Manifest Variables on Estimates of Statistical Power?

In a second step, we further examined the effect of the construct reliability on statistical power by investigating the loading strength of the manifest variables. To this end we conducted a second simulation following the steps outlined above while varying the sample size from 50 to 200 and the loadings of the manifest variables from 0.5 to 0.9. Figure 3B displays the obtained power curves, illustrating statistical power dependent on total sample size and loading strength of the manifest variables. For this estimate, the factor inter-correlation was set to 0.3. Again, increasing sample size lead to higher statistical power. However, the loading strength impacted estimates of statistical power: the higher the loadings of the manifest variables, the higher the statistical power achieved by a given sample size. For reasonably high levels of factor loadings (≥ 0.7) a statistical power of 95% was reached for $N = 120$.

3.3 How can Simulations Inform us About Type-I Error Rates?

The main part of this tutorial has focused on estimating statistical power that is required to reject a null hypothesis when it is indeed false (Type-II error). However, it is necessary to also consider the reverse — rejecting a null hypothesis that is in fact true (Type-I error). In our example, Type-I error would be committed by misidentifying Model 3 as the best fitting model when in reality, the data were generated from either Models 1 or 2. We can leverage the simulation approach to quantify Type-I error rates dependent on the different simulation conditions (e.g., inter-factor correlations, factor loadings). We conducted an analogous analysis to the one outlined above (sections 2.1.1. – 2.1.6), with the difference that Models 1 and 2 are defined as the generating population models (step 3, section 2.1.3). The outcomes of this analysis show the probability of erroneously identifying an overly complex model, when the more simplistic models generated the data (see Figure 4).

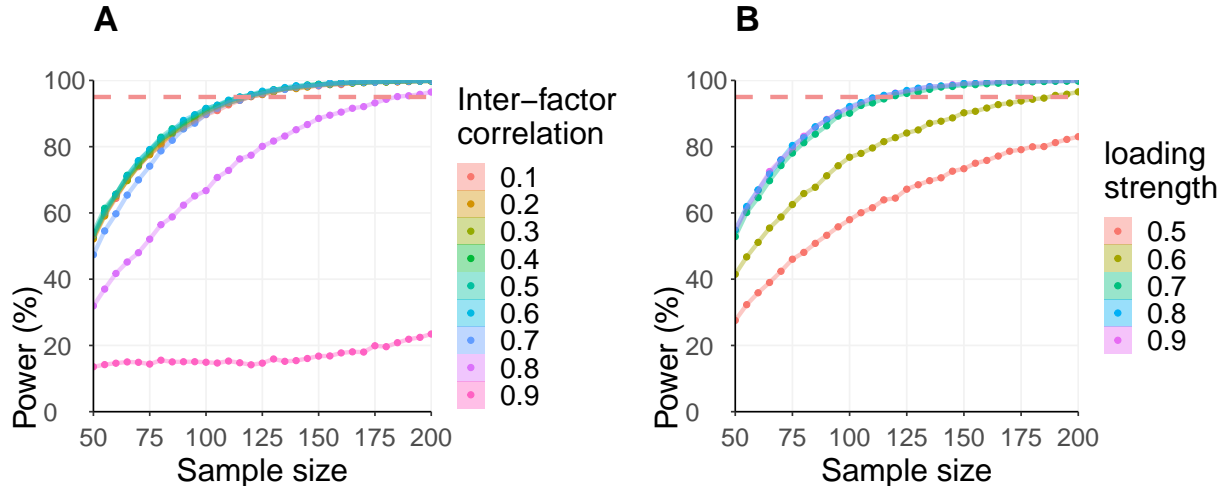


Figure 3. Simulation results: Estimated statistical power for recovering the data-generating model, i.e. Model 3 (y-axis) as a function of sample size (x-axis) with varying the factor inter-correlations when the loading strength was fixed to 0.7 (A) ; and with varying the loading strength when the factor inter-correlation was fixed to 0.3 (B). For each combination of parameters, 10,000 datasets were simulated. The red dashed horizontal line positioned at 95% discovery rate reflects the cut-off for identifying the appropriate sample size. The syntax to generate these plots with `ggplot()` can be found in the supplementary R Code on GitHub

The results of this analysis reveal that even for low levels of loadings strength (< 0.7), the probability of erroneously picking Model 3 as the best fitting model is reasonably low ($< 7\%$). Likewise, this probability remains reasonably low for different levels of inter-factor correlations for all sample sizes ($< 5\%$). Thus, in this specific model comparison Type-I error probability does not appear to be the main aspect driving sample size requirements for an informative research design.

3.4 Deriving a Sample Estimate From the Simulation Results

To derive a joint sample size estimate from the different parameters in the simulation, it is crucial to reiterate the theoretical assumptions on the strength of all parameters investigated in the simulation. In our illustration, it is reasonable to assume that the latent factors in the tri-partite model (Model 3, Figure 1) are only moderately correlated (below

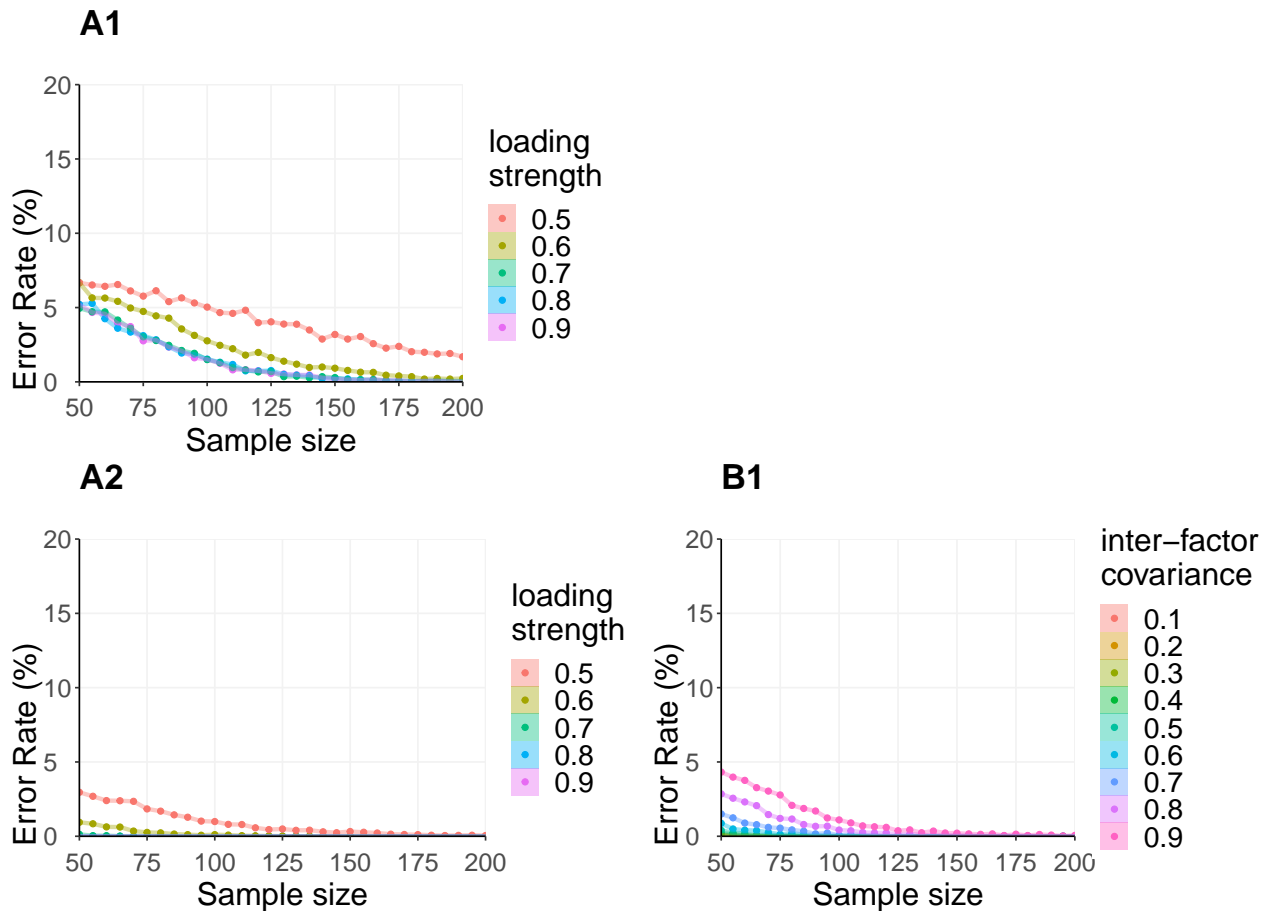


Figure 4. Simulation results: Probability to erroneously pick Model 3 as the best fitting model when indeed Model 1 (A1) or Model 2 (A2/B2) generated the data, as a function of total sample size (x-axis) and the loadings strength of the manifest variables (A1/A2) or the strength of the covariance of the latent factors (B2). Note that the y-axis reaches from 0 to 20% for the sake of visibility.

0.7). Therefore, our model comparison requires a sample size of $N = 120$ to reach statistical power of 95% to adjudicate between the theoretical models, if the hypothesized model (Model 3) generates the data. We further assume that the manifest variables capture the latent constructs reasonably well — that is, exerting a loading strength of 0.7 or above. The results of the simulation suggest that $N = 120$ would suffice to achieve 95% with regard to the reliability of the manifest factors. In our example, the results of both analyses align, such that a sample size of $N = 120$ suffices to achieve a power of 95% regarding both the separability of the theoretical constructs and the reliability of the indicators. In cases where the derived sample estimate differs between the two analyses we suggest to derive the more conservative estimate of required sample size from the two simulations. The results from the Type-I error analysis further indicate, that the probability of erroneously choosing an overly complex model is generally low. Therefore, the probability of committing a Type-I error for the sample size derived from the simulations on model separability and measurement reliability ($N = 120$) is negligible.

4. Discussion and Conclusion

SEM offers a powerful tool for modeling multivariate relationships between psychological constructs. Maximizing its utility in psychological sciences hinges on the integration of its complex characteristics with power estimation – another key tenet of scientific rigor. Our work exemplifies a roadmap for navigating some of the challenges in harmonizing these two efforts by implementing a-priori power estimation within an SEM framework. In the case of information-theoretic model comparisons for non-nested models, statistical power cannot be determined analytically, i.e., model comparisons cannot be performed by using χ^2 -based statistics. Monte Carlo simulations provide a useful tool to circumvent this methodological challenge (Muthén & Muthén, 2002), allowing researchers to derive informed decisions on sample size planning to ensure sufficient statistical power and prevent uninformative or misleading research (Button et al., 2013). In this study, we

illustrate a step-by-step approach to implementing a Monte Carlo simulation in the statistical programming language R to estimate the required sample size for model comparisons of non-nested models. The complete code is available online to yield a starting point for researchers who want to implement their own Monte Carlo power simulations (<https://github.com/ebuchberger/Estimating-Power-for-SEM>).

The results from our illustration highlight the importance of two model parameters: the separability of the theoretical models on the basis of factor inter-correlation, and the reliability with which the construct can be measured on the basis of the loading strength of the manifest variables. These results align with theoretical knowledge on SEM Wolf et al. (2013), but they provide additional insights on the tangible impact of these parameters on statistical power estimation in the specific research design discussed here and in other SEMs more broadly.

First, we showed that detecting differences in model fit requires increasingly many participants the higher the inter-correlation of the factors turns out practically. Therefore, it is warranted to consider the actual separability of the constructs in question to avoid underpowered sampling plans. We acknowledge that committing to specific parameter values, such as factor inter-correlations, can be difficult when setting up the population models for the simulations. Thus, we recommend researchers to consult existing literature in order to make informed assumptions on the relationship among different latent factors. For cases in which no or little evidence exists on the relationship between the constructs under investigation, we recommend to enter these inter-correlations into the simulation as parameter of interest (as demonstrated in the example of this paper) and to then carefully evaluate what level of resolution is required when separating inter-related factors depending on practical relevance.

Second, the results on the loadings of manifest variables stresses the importance of employing indicators that capture the latent construct with high reliability (Zuo, Xu, &

Milham, 2019). Considering the issue of reliability is a crucial step in experimental design, as the reliability of the indicators heavily influences the practical ability to identify the targeted latent construct with sufficient precision and differentiate between competing models. Again, committing to population values can be difficult, especially if there is no prior work from which to draw values. In this case, we still advice to estimate plausible values from related scientific areas. For instance, selecting factor loading strength in the tri-partite model is not straightforward. However, we can do a literature search on similar models and take the reported test- reliability (e.g., a test-retest) or construct reliability (e.g., coefficient alpha) and translate this back into factor loadings strength and residual loadings (see Brandmaier et al., 2018 for examples on how to parametrize longitudinal models from minimal available information). If information about test reliability is available, we set the loadings to the square-root of the reliability to take this prior information into account. Then, this yields the following construct reliability for our design (assuming identical loadings across k items):

$$\alpha = \frac{\sum_k \lambda^2}{\sum_k \lambda^2 + 1 - \lambda^2} \quad (1)$$

If instead, prior knowledge about construct reliability exists and this reliability estimate α was based on k items, we solve Equation 1 for λ and obtain an estimate that may serve to inform our factor loadings:

$$\lambda = \sqrt{\frac{\alpha}{\alpha + k - \alpha \cdot k}} \quad (2)$$

For example, taking from a literature search that the construct reliability of the construct of interest (reported as Cronbach's alpha) was 0.7 and the estimate was based on $k=4$ tests, we set our loadings to $\lambda = 0.61$ (using Eq.2). If a study reported a test-retest correlation of a single test to be $r=0.8$, then we set the factor loading to $\lambda = \sqrt{0.8} = 0.64$. If in doubt, we suggest to use conservative values rather than too liberal values to avoid

underpowered studies.

While we demonstrate in this tutorial that simulation-based techniques help overcome methodological shortcomings from traditional analytical approaches, the randomness that is inherent to such approaches may seem at odds with efforts in results reproducibility. Here, a clear-cut solution is through transparent and accessible documentation of the simulation code and adequate software management. That is, researchers should make use of repositories on platforms such as GitHub or the Open Science Framework to make code scripts available to other researchers, and clearly specify the versions of software programs and packages that were used (e.g. via containerization or at least in written form). We thus urge researchers to seriously consider the issue of reproducibility for their simulation analyses, to make outcomes of such a-priori power estimates easily accessible to other researchers and reviewers (for a more detailed guideline on a reproducible workflow in R see Peikert, Lissa, & Brandmaier, 2021).

In sum, we argue that sample size planning in the case of non-nested SEM comparisons requires a simulation-based approach. In particular, we highlighted that the separability of theoretical constructs, as well as the reliability of the measures, have a major impact on estimates of statistical power. To this end we hope that this tutorial advances the use of simulation-based approaches to estimating statistical power in cases where analytical approaches do not suffice.

References

- Akaike, H. (1974). A new look at the statistical model identification. *IEEE Transactions on Automatic Control*, 19(6), 716–723. <https://doi.org/10.1109/TAC.1974.1100705>
- Anderson, S. F., & Maxwell, S. E. (2017). Addressing the “replication crisis”: Using original studies to design replication studies with appropriate statistical power. *Multivariate Behavioral Research*, 52(3), 305–324. <https://doi.org/10.1080/00273171.2017.1289361>
- Baltes, P. B., Reese, H. W., & Nesselroade, J. R. (1988). *Life-span developmental psychology: Introduction to research methods*. Hillsdale, NJ: Erlbaum.
- Bentler, P. M. (1990). Comparative fit indexes in structural models. *Psychological Bulletin*, 107(2), 238. <https://doi.org/10.1037/0033-2909.107.2.238>
- Bentler, P. M., & Bonett, D. G. (1980). Significance tests and goodness of fit in the analysis of covariance structures. *Psychological Bulletin*, 88(3), 588.
- Bentler, P. M., & Chou, C. (1987). Practical issues in structural modeling. *Sociological Methods & Research*, 16(1), 78–117.
- Bollen, K. A. (1989). *Structural equations with latent variables* (Vol. 210). John Wiley & Sons.
- Boomsma, A. (1985). Nonconvergence, improper solutions, and starting values in LISREL maximum likelihood estimation. *Psychometrika*, 50(2), 229–242.
- Brandmaier, A. M. (2020). Optimal planned missing data design for linear latent growth curve models. *Behavior Research Methods*, 52, 1445–1458. <https://doi.org/10.3758/s13428-019-01325-y>
- Brandmaier, A. M., Oertzen, T. von, Ghisletta, P., Hertzog, C., & Lindenberger, U. (2015). LIFESPAN: A tool for the computer-aided design of longitudinal studies. *Frontiers in Psychology*, 6. Retrieved from <https://www.frontiersin.org/article/10.3389/fpsyg.2015.00272>
- Brandmaier, A. M., Oertzen, T. von, Ghisletta, P., Lindenberger, U., & Hertzog, C. (2018). Precision , reliability, and effect size of slope variance in latent growth curve models:

- Implications for statistical power analysis. *Frontiers in Psychology*, 9.
<https://doi.org/10.3389/fpsyg.2018.00294>
- Buchberger, E. S., Brandmaier, A. M., Lindenberger, U., Werkle-Bergner, M., & Ngo, C. T. (2022). *The process structure of memory abilities in early and middle childhood*.
- Button, K. S., Ioannidis, J. P., Mokrysz, C., Nosek, B. A., Flint, J., Robinson, E. S., & Munafò, M. R. (2013). Power failure: Why small sample size undermines the reliability of neuroscience. *Nature Reviews Neuroscience*, 14(5), 365–376.
- Canada, K. L., Hancock, G. R., & Riggins, T. (2021). Developmental changes in episodic memory across early- to mid-childhood: Insights from a latent longitudinal approach. *Memory*, 30(3), 1–14. <https://doi.org/10.1080/09658211.2021.2006233>
- Cattell, R. B. (1971). *Abilities: Their structure, growth, and action*.
- Chalmers, R. P., & Adkins, M. C. (2020). Writing effective and reliable monte carlo simulations with the SimDesign package. *The Quantitative Methods for Psychology*, 16(4), 248–280.
- Cohen, J. (1962). The statistical power of abnormal-social psychological research: A review. *The Journal of Abnormal and Social Psychology*, 65(3), 145.
<https://doi.org/10.1037/h0045186>
- Cohen, J. (1988). *Statistical power analysis for the behavioral sciences*. Hillsdale: NJ Lawrence Earlbaum Associates.
- Efron, B., & Tibshirani, R. J. (1994). *An introduction to the bootstrap*. CRC press.
- Hedden, T., & Yoon, C. (2006). Individual differences in executive processing predict susceptibility to interference in verbal working memory. *Neuropsychology*, 20(5), 511–528.
<https://doi.org/10.1037/0894-4105.20.5.511>
- Henson, R. N., Campbell, K. L., Davis, S. W., Taylor, J. R., Emery, T., Erzinclioglu, S., & Kievit, R. A. (2016). Multiple determinants of lifespan memory differences. *Scientific Reports*, 6(1), 32527. <https://doi.org/10.1038/srep32527>
- Horn, J. L. (1970). Organization of data on life-span development of human abilities. In

- Life-span developmental psychology* (pp. 423–466). Elsevier.
- Horn, J. L. (1978). Human ability systems. *Life-Span Development and Behavior*.
- Hu, L., & Bentler, P. M. (1999). Cutoff criteria for fit indexes in covariance structure analysis: Conventional criteria versus new alternatives. *Structural Equation Modeling: A Multidisciplinary Journal*, 6(1), 1–55. <https://doi.org/10.1080/10705519909540118>
- Jak, S., Jorgensen, T. D., Verdam, M. G. E., Oort, F. J., & Elffers, L. (2020). Analytical power calculations for structural equation modeling: A tutorial and shiny app. *Behavior Research Methods*. <https://doi.org/10.3758/s13428-020-01479-0>
- Jobst, L. J., Bader, M., & Moshagen, M. (2021). A tutorial on assessing statistical power and determining sample size for structural equation models. *Psychological Methods*. <https://doi.org/10.1037/met0000423>
- Jöreskog, K. G., & Goldberger, A. S. (1975). Estimation of a model with multiple indicators and multiple causes of a single latent variable. *Journal of the American Statistical Association*, 70(351), 631–639. <https://doi.org/10.1080/01621459.1975.10482485>
- Kievit, R. A., Brandmaier, A. M., Ziegler, G., Harmelen, A.-L. van, Mooij, S. M. M. de, Moutoussis, M., ... Dolan, R. J. (2018). Developmental cognitive neuroscience using latent change score models: A tutorial and applications. *Developmental Cognitive Neuroscience*, 33, 99–117. <https://doi.org/10.1016/j.dcn.2017.11.007>
- MacCallum, R. C., & Austin, J. T. (2000). Applications of structural equation modeling in psychological research. *Annual Review of Psychology*, 51(1), 201–226. <https://doi.org/10.1146/annurev.psych.51.1.201>
- Marr, D. (1971). Simple memory: A theory for archicortex. *Philos. Trans. R. Soc. London*, 262, 23–81.
- Maxwell, S. E. (2004). The persistence of underpowered studies in psychological research: Causes, consequences, and remedies. *Psychological Methods*, 9(2), 147. <https://doi.org/10.1037/1082-989X.9.2.147>
- McClelland, J. L., McNaughton, B. L., & O'Reilly, R. C. (1995). Why there are

- complementary learning systems in the hippocampus and neocortex: Insights from the successes and failures of connectionist models of learning and memory. *Psychological Review*, 102(3), 419–457. <https://doi.org/10.1037/0033-295X.102.3.419>
- McCrae, R., & Costa, P. (1985). Updating norman's "adequate taxonomy": Intelligence and personality dimensions in natural language and in questionnaires. *Journal of Personality and Social Psychology*, 49, 710–721. <https://doi.org/10.1037//0022-3514.49.3.710>
- Metropolis, N., Rosenbluth, A. W., Rosenbluth, M. N., Teller, A. H., & Teller, E. (1953). Equation of state calculations by fast computing machines. *The Journal of Chemical Physics*, 21(6), 1087–1092. <https://doi.org/10.1063/1.1699114>
- Miller, M. R., Giesbrecht, G. F., Müller, U., McInerney, R. J., & Kerns, K. A. (2012). A latent variable approach to determining the structure of executive function in preschool children. *Journal of Cognition and Development*, 13(3), 395–423. <https://doi.org/10.1080/15248372.2011.585478>
- Mirsky, A. F., Anthony, B. J., Duncan, C. C., Ahearn, M. B., & Kellam, S. G. (1991). Analysis of the elements of attention: A neuropsychological approach. *Neuropsychology Review*, 2(2), 109–145. <https://doi.org/10.1007/BF01109051>
- Miyake, A., Friedman, N. P., Emerson, M. J., Witzki, A. H., Howerter, A., & Wager, T. D. (2000). The unity and diversity of executive functions and their contributions to complex “frontal lobe” tasks: A latent variable analysis. *Cognitive Psychology*, 41(1), 49–100. <https://doi.org/10.1006/cogp.1999.0734>
- Muthén, L. K., & Muthén, B. O. (2002). How to use a monte carlo study to decide on sample size and determine power. *Structural Equation Modeling: A Multidisciplinary Journal*, 9(4), 599–620. https://doi.org/10.1207/S15328007SEM0904_8
- Neale, M. C. (2000). Individual fit, heterogeneity, and missing data in multigroup sem. *Modeling Longitudinal and Multiple-Group Data: Practical Issues, Applied Approaches, and Specific Examples*. Hillsdale, NJ: Lawrence Erlbaum Associates.
- Norman, K. A., & O'Reilly, R. C. (2003). Modeling hippocampal and neocortical

- contributions to recognition memory: A complementary-learning-systems approach. *Psychological Review*, 110(4), 611–646. <https://doi.org/10.1037/0033-295X.110.4.611>
- Nyberg, L. (1994). A structural equation modeling approach to the multiple memory systems question. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 20(2), 485–491. <https://doi.org/10.1037/0278-7393.20.2.485>
- Paxton, P., Curran, P. J., Bollen, K. A., Kirby, J., & Chen, F. (2001). Monte carlo experiments: Design and implementation. *Structural Equation Modeling: A Multidisciplinary Journal*, 8(2), 287–312. https://doi.org/10.1207/S15328007SEM0802_7
- Peikert, A., Lissa, C. J. van, & Brandmaier, A. M. (2021). Reproducible research in r: A tutorial on how to do the same thing more than once. *Psych*, 3(4), 836–867. <https://doi.org/10.3390/psych3040053>
- Pornprasertmanit, S., Miller, P., Schoemann, A., & Jorgensen, T. D. (2021). Simsem: SIMulated structural equation modeling. *R Package Version 0.5-16*. Retrieved from <https://CRAN.R-project.org/package=simsem>
- Raftery, A. E. (1986). Choosing models for cross-classifications. *American Sociological Review*, 51(1), 145–146.
- Rolls, E. T. (2016). Pattern separation, completion, and categorisation in the hippocampus and neocortex. *Neurobiology of Learning and Memory*, 129, 4–28. <https://doi.org/10.1016/j.nlm.2015.07.008>
- Rosseel, Y. (2012). Lavaan: An r package for structural equation modeling and more version 0.5-12 (BETA). *Journal of Statistical Software*, 37. <https://doi.org/10.18637/jss.v048.i02>
- Rossi, J. S. (1990). Statistical power of psychological research: What have we gained in 20 years? *Journal of Consulting and Clinical Psychology*, 58(5), 646–656. <https://doi.org/10.1037/0022-006X.58.5.646>
- Satorra, A., & Saris, W. E. (1985). Power of the likelihood ratio test in covariance structure analysis. *Psychometrika*, 50, 83–90. <https://doi.org/10.1007/BF02294150>
- Schapiro, A. C., Turk-Browne, N. B., Botvinick, M. M., & Norman, K. A. (2017).

- Complementary learning systems within the hippocampus: A neural network modelling approach to reconciling episodic memory with statistical learning. *Philosophical Transactions of the Royal Society B: Biological Sciences*, 372(1711), 20160049.
<https://doi.org/10.1098/rstb.2016.0049>
- Schoemann, A. M., Miller, P., Pornprasertmanit, S., & Wu, W. (2014). Using monte carlo simulations to determine power and sample size for planned missing designs. *International Journal of Behavioral Development*, 38(5), 471–479.
<https://doi.org/10.1177/0165025413515169>
- Schwarz, G. (1978). Estimating the dimension of a model. *The Annals of Statistics*, 6(2), 461–464.
- Sedlmeier, P., & Gigerenzer, G. (1989). Do studies of statistical power have an effect on the power of studies? *Psychological Bulletin*, 105(2), 309.
<https://doi.org/10.1037/0033-2909.105.2.309>
- Spearman, C. (1904). "General intelligence" objectively determined and measured. *American Journal of Psychology*, 15, 201–293.
- Squire, L. R. (1987). *Memory and brain*. New York, NY, US: Oxford University Press.
- Steiger, J. H. (2016). Notes on the steiger–lind (1980) handout. *Structural Equation Modeling: A Multidisciplinary Journal*, 23(6), 777–781.
<https://doi.org/10.1080/10705511.2016.1217487>
- Tucker-Drob, E. M., Brandmaier, A. M., & Lindenberger, U. (2019). Coupled cognitive changes in adulthood: A meta-analysis. *Psychological Bulletin*, 145(3), 273–301.
<https://doi.org/10.1037/bul0000179>
- Tulving, E. (1972). Episodic and semantic memory. *Organization of Memory*, 1, 381–403.
- Vankov, I., Bowers, J., & Munafò, M. R. (2014). Article commentary: On the persistence of low power in psychological science. *Quarterly Journal of Experimental Psychology*, 67(5), 1037–1040. <https://doi.org/10.1080/17470218.2014.885986>
- Wang, Y. A., & Rhemtulla, M. (2021). Power analysis for parameter estimation in structural

equation modeling: A discussion and tutorial. *Advances in Methods and Practices in Psychological Science*, 4(1), 2515245920918253.

Wolf, E. J., Harrington, K. M., Clark, S. L., & Miller, M. W. (2013). Sample size requirements for structural equation models: An evaluation of power, bias, and solution propriety. *Educational and Psychological Measurement*, 73(6), 913–934.
<https://doi.org/10.1177/0013164413495237>

Zuo, X.-N., Xu, T., & Milham, M. P. (2019). Harnessing reliability for neuroscience research. *Nature Human Behaviour*, 3(8), 768–771. <https://doi.org/10.1038/s41562-019-0655-x>